

HOMEWORK #1 REPORT

Granny's Kitchen

Enrico Dalla Mora, Rocco Scarano, Nicolò Chillè, and Federico Caroppo

Group Name: *La Lobby* - Group ID: 3

Abstract

The purpose of this document is to describe the process of creating *foley sounds* by means of granular synthesis in the **SuperCollider** environment and implementing an integrated graphical user interface in order to allow, to some extent, the customization of the sound parameters by the user. Given that foley sounds are usually employed in audiovisual content as a mean to enrich the soundscape of a scene, the sounds herein analyzed were developed to mimic the sounds that are usually perceived in a specific environment.

Keywords: granular, synthesis, foley

The synthesis process started with the selection of the environment to which the synthesized sounds would belong. The choice fell on an ordinary life soundscape: the one of a *kitchen*. Through granular synthesis, the sounds associated to the following objects were recreated:

- a **gas stove** being turned on - ①
- an **onion** being cut by a knife cutting board - ②
- two **glasses** clashing - ③
- boiling **kettle** - ④

Even though granular synthesis is, to an extent, affected by the micro-sampled track, the choice of said track was solely based on the audio context and fell on "*Le Festin*" by Camille and Michael Giacchino, a song from the soundtrack of the hugely popular animated movie "*Ratatouille*" (2007), which is set in a kitchen for the most part.

The first section of this report will focus on the approach adopted to achieve the synthesis of the aforementioned sounds in the SuperCollider coding environment.

Note

Some of the UGen input values have been parameterized in order to allow the user to shape the sound via the GUI (i.e. the corresponding variable value will depend on the related slider/knob). The parameters, reported in the code snippets, are the following: *pos*, *rat*, *tickSpeed*, *len* and *dens*.

1. Granular Synthesis and SuperCollider

The sound synthesis method known as **granular synthesis** is based on the production of a high density of small acoustic events (the so-called 'grains') that are typically under 100 ms in duration, obtained via the *microsampling* of a given audio signal. Adding a large number of grains together and organizing them will result in the synthesis of a texture, therefore generating a sound from a macroscopic standpoint.

SuperCollider (which will be referred to, from now on, as SC) provides a UGen, called `GrainBuf`, that implements granular synthesis starting from the audio signal stored in a buffer (~b3). The class is structured as follows:

Code 1. The `GrainBuf` class.

```
1 GrainBuf.ar(numChannels: 1, trigger: 0, dur: 1, sndbuf, rate:
    1, pos: 0, interp: 2, pan: pan, envbufnum: -1, maxGrains:
    512, mul: 1, add: 0);
```

The most significant parameters, i.e. the ones that were involved in the development of the sounds, are:

- **trigger**: the **trigger** to start a new grain (either at audio rate or at control rate). Two kinds of trigger were used, namely:
 - the *impulse oscillator* `Impulse`, which is more suited for impulsive audio events such as sounds ② and ③;
 - the *random impulses generator* `Dust`, more fitting for rough, noisy sounds such as ① and ④.
- **dur**: indicates the temporal **size** of the grain (in *seconds*);
- **rate**: the **playback rate** of the sampled sound, affecting pitch and timbre;
- **pos**: the **position** for the grain to start with, relative to the input sound (range [0, 1] - where 0 represents the beginning of the sample and 1 corresponds with its end).

The sound synthesized by each `GrainBuf` object class was further shaped by an **envelope** and, possibly, by effects (e.g. *reverb* or *resonance*). In particular, the purpose of the reverb is to enhance the spatial details linked to the perception of the sounds in a real room.

2. Implementation

In the present section, each one of the foley sounds will be discussed individually and in detail. We chose to represent the basic synthesis chain for readability purposes. Actually this elementary component is combined and developed in many different ways throughout the various sounds. The actual UGen graphs can be found inside the folder "UGenGraphs" on the GitHub page.

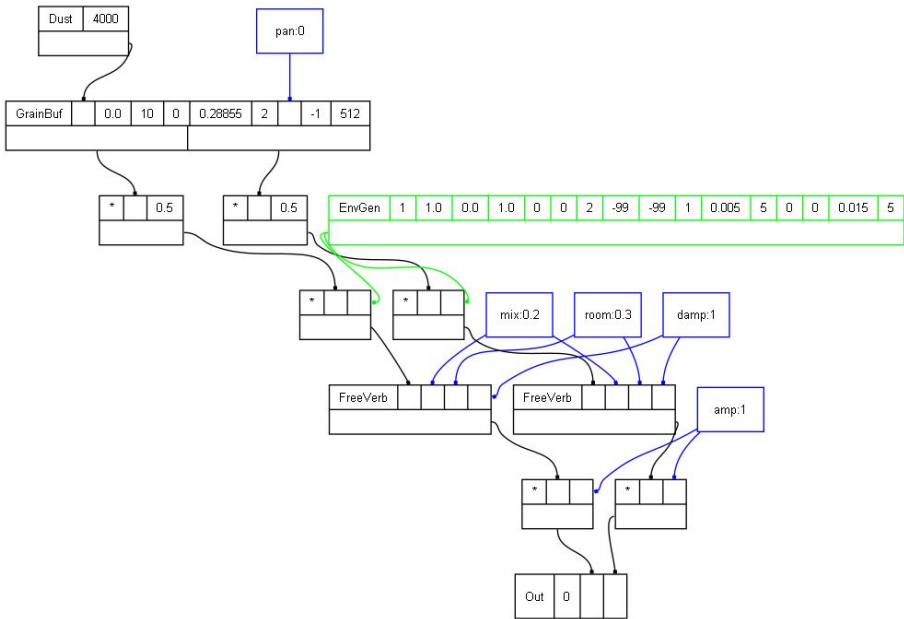


Figure 1. Simplified UGen graph.

2.1 Gas Stove ①

The objective was to reproduce the complex sound of a gas stove being turned on, from the gas flow rustling to the flame rumble, including the regular pulsing of the spark module. Thus, the synthesized signal is structured in four superimposed components, namely the gas flow, the ticking of the sparks, the burst of the lighting flame and the flames rumble. The transition between said components is managed via a number of envelopes.

2.1.1 Signal 1 - Gas flow

The first component of the final signal corresponds to the *rustling of the gas flow*. Being this a **noise-like sound**, the trigger signal was generated using the `Dust` UGen. The average number of impulses (density) given as input is very high, and thus the duration of the grains must be low to avoid problems linked to overlapping. As said before, the *rate* parameter affects the perceived pitch. *No envelope* was used to further shape the sound.

Code 2. The gas flow signal.

```
1 sig1 = GrainBuf.ar(numChannels: 2, trigger: Dust.ar(1469.89), dur: 0.04912,
  sndbuf: ~b3, rate: 1.89+((rat*2)-1), pos: 0.00000, interp: 2, pan: pan,
  envbufnum: -1, maxGrains: 512);
```

2.1.2 Signal 2 - Spark module

The base signal is still noise-like, but it is characterized by a much higher density of grains in order to achieve a *richer high-frequency spectrum*. The pulsing effect is achieved via an **impulse-shaped periodic envelope** (Figure 2) applied to the signal amplitude.

Code 3. The ticking sparks signal and envelope.

```
1 sig2 = GrainBuf.ar(numChannels: 2, trigger: Dust.ar(4000), dur: 0.00000, sndbuf:
  ~b3, rate: 0, pos: 0.28855, interp: 2, pan: pan, envbufnum: -1, maxGrains:
  512);
2 env2 = EnvGen.ar(Env([0, 1, 0], [0.005, 0.015], curve: 0), Impulse.kr(4*
  tickSpeed));
```




Figure 2. Spark ticking amplitude envelope.

2.1.3 Signal 3 - Flame rumble

The third signal represents the flame rumble that follows the flame lighting burst (i.e. signal 4). The peculiarity of this sound is the way the transition between the spark ticking and the ramble is handled: the envelope named `env3` (Figure 3-a) is a `LFPulse` characterized by a very low frequency and a long duty cycle, and it handles the ignition of the flame in correspondence with the last spark, while `env31` is an inverted square envelope (refer to Code 4) with respect to `env3` and handles the deactivation of the spark ticking after the igniting one. The duty cycle and the initial phase are slightly different from the `env3` ones in order to allow the last spark (the igniting one) to be heard.

Code 4. The flame rumble signal and envelopes.

```
1 sig3 = GrainBuf.ar(numChannels: 2, trigger: Dust.ar(388), dur: 0.00391, sndbuf:
  ~b3, rate: 0.29+((rat*0.4)-0.2), pos: 0.16565+((pos*2)-1), interp: 2, pan:
  pan, envbufnum: -1, maxGrains: 512);
2 env3 = LFPulse.ar(0.08, 0.8, 0.8);
3 env31= 1 - LFPulse.ar(0.08, 0.79, 0.79);
```




Figure 3. env3 and env31 amplitude envelopes.

2.1.4 Signal 4 - Ignition burst

Code 5. The flame burst signal and envelope.

```
1 sig4=GrainBuf.ar(numChannels: 2, trigger: Dust.ar(388), dur: 0.01291, sndbuf: ~b3
  , rate: 0.30, pos: 0.16565+((pos*2)-1), interp: 2, pan: pan, envbufnum: -1,
  maxGrains: 512);
2 env4=EnvGen.ar(Env([1, 0], [0.7], curve: 0), env3);
```

The last signal is also a rough sound, and thus employs *Dust* as a trigger signal generator, but its input density is much lower than the flame rumble one. This is due to the necessity of a lower pitch to mimic the fire blaze. The audio event must be impulsive, as the overall signal transitions into the lighted stove rumble: this effect is achieved by means of a single decaying envelope (for code reference, see Code 5 and Figure 3).




Figure 4. Blaze amplitude envelope.

2.1.5 Final sound - Igniting stove

To generate the complete sound, the four signals are amplified (via four amplitude factors), modulated by means of the dedicated envelopes and added together to be run through a simple reverb and sent out as a stereo output.

Code 6. The flame rumble signal and envelopes.

```

1 sig = (sig1*amp1) + (sig2*amp2*env2*env31) + (sig3*amp3*env3) + (sig4*amp4*env4*
  env3);
2 Out.ar(busNum, FreeVerb.ar(sig, mix, room, damp)*amp);

```

2.2 Onions ②

The sound in question is comprised of *two components* (signals): the knife cutting through an hypothetical vegetable – sig1 – and the dull thud of the knife hitting the cutting board – sig2. Each of this sound components is of **periodic and impulsive character**, thus the chosen trigger signal generator is **Impulse**. The trigger of the board thud is, obviously, slightly delayed with respect to the one relative to the first signal. The overall signal is then amplified and sent in output through a reverb.

Code 7. The cutting knife signals and envelope.

```

1 sig1 = GrainBuf.ar(numChannels: 2, trigger: Impulse.ar(0.5+tickSpeed), dur: 0.4,
  sndbuf: ~b3, rate: 0.5+((rat*0.4)-0.2), pos: 0.27, interp: 2, pan: pan,
  envbufnum: -1, maxGrains: 512);
2 sig2 = GrainBuf.ar(numChannels: 2, trigger: Impulse.ar(0.5+tickSpeed, 0.87), dur:
  0.02524, sndbuf: ~b3, rate: 0.23+((rat*0.2)-0.1), pos: 0.95064-(pos*0.2),
  interp: 2, pan: pan, envbufnum: -1, maxGrains: 512);
3 sig = (sig1*amp1) + (sig2*amp2);
4 Out.ar(busNum+2, FreeVerb.ar(sig, mix, room, damp)*amp);

```

2.3 Glasses clashing ③

The glass clinking effect was achieved by *feeding a single, very short grain to a resonator* UGen: **DynKlank**. The purpose of said UGen is to simulate the resonant modes of an object: each mode is given a ring time (corresponding to the time necessary for the mode to decay by 60 dB). The signal's amplitude is shaped through one envelope, which gives the impulsive quality and is then fed to a reverb. The parameters **pitch** and **dur** are obtained starting from **rat** and **len**. Differently from the other sounds, this one is not played in loop, but it is activated every time the user hits the play button. For this reason, the **SynthDef** is instantiated every time the sound is played and therefore there is a final envelope (**env**) which is used to free the synth after the resonator stops playing.

Code 8. The clashing glasses signal and envelopes.

```

1 sig = GrainBuf.ar(numChannels: 2, trigger: Impulse.ar(319.09), dur: 0.09134,
  sndbuf: ~b3, rate: 1.61, pos: 0.17028+((pos*0.05)-0.025), interp: 2, pan: 0,
  envbufnum: -1, maxGrains: 512);
2 env1 = EnvGen.ar(Env([1, 0], [0.03], curve: [500]), 1, 1, 0, 1, 0);
3 sig = DynKlank.ar(`[[800*1.2*pitch, 2684*1.2*pitch, 4612*1.2*pitch, 6892*1.2*
  pitch], [1, 0.25, 0.165, 0.0625], [5+dur, 4+(dur/2), 3+(dur/3), 2+(dur/4)]]),
  sig*env1*0.1);
4 rev = FreeVerb.ar(sig, mix, room, damp);
5 panner = Pan2.ar(rev*amp, pan);
6 env = EnvGen.ar(Env([1, 0], [6+length], curve: [-5]), 1, 1, 0, 1, 2);
7 Out.ar(busNum+(2*2), panner);

```

2.4 Boiling kettle ④

The idea behind the boiling water sound is to smooth a superimposition of four different rough, noisy signals. For each one of them, a *Dust* trigger signal generator was employed (using various input densities) and a **very small duration** of the grains was set.

Code 9. The boiling kettle signal and envelopes.

```

1 sig1 = GrainBuf.ar(numChannels: 2, trigger: Dust.ar(20+((dens*40)-20)), dur:
    0.01751, sndbuf: ~b3, rate: 1+((rat*1)-0.5), pos: 0.06094+((pos*0.1)-0.05),
    interp: 2, pan: pan, envbufnum: -1, maxGrains: 512);
2 sig2 = GrainBuf.ar(numChannels: 2, trigger: Dust.ar(10+((dens*20)-10)), dur:
    0.01751, sndbuf: ~b3, rate: 2+((rat*2)-1), pos: 0.16094+((pos*0.2)-0.1),
    interp: 2, pan: pan, envbufnum: -1, maxGrains: 512);
3 sig3 = GrainBuf.ar(numChannels: 2, trigger: Dust.ar(30+((dens*60)-30)), dur:
    0.01751, sndbuf: ~b3, rate: 3+((rat*3)-1.5), pos: 0.26094+((pos*0.2)-0.1),
    interp: 2, pan: pan, envbufnum: -1, maxGrains: 512);
4 sig4 = GrainBuf.ar(numChannels: 2, trigger: Dust.ar(855), dur: 0.01234, sndbuf:
    ~b3, rate: 9.61+((rat*3)-1.5), pos: 0.31664, interp: 2, pan: pan, envbufnum:
    -1, maxGrains: 512);
5
6 sig = (sig1*amp1) + (sig2*amp2) + (sig3*amp3) + (sig4*amp4);
7 Out.ar(busNum+(3*2), FreeVerb.ar(sig, mix, room, damp)*amp*0.25);

```

2.5 Scoping

In order to show the the sounds both in the time and frequency domain the sounds have been implemented as to output on sequential busses (4, 6, 8, 10) which are then routed to the scopes and to the final output bus (0) by the *scoping* SynthDef.

Code 10. Scoping SynthDefs and instantiation

```

1 SynthDef(\scoping, {
2   arg bus, buff;
3   ScopeOut2.ar(inputArray: In.ar(bus,2), scopeNum: buff, maxFrames: 8092,
    scopeFrames: 1024);
4   Out.ar([0,1], In.ar(bus, 2));
5   }).add;
6
7 {\scope1 = Synth(\scoping, [\buff, f1.bufnum, \bus, busNum])}.defer(0.1);
8 {\scope2 = Synth(\scoping, [\buff, f2.bufnum, \bus, busNum+2])}.defer(0.1);
9 {\scope3 = Synth(\scoping, [\buff, f3.bufnum, \bus, busNum+(2*2)])}.defer(0.1);
10 {\scope4 = Synth(\scoping, [\buff, f4.bufnum, \bus, busNum+(3*2)])}.defer(0.1);

```

The *ScopeOut2* UGen fills a buffer, which has previously been allocated, with the scoping data in the time domain. This buffer will then be plotted by the UGen *ScopeWiew* in the GUI implementation. For what concerns the frequency scope it's easier to view the spectrum since the *FreqScopeWiew* UGen provides the *inBus()* method.

3. Graphical User Interface (GUI)

The GUI that has been designed for the assignment is shown in Figure 4. The GUI is divided in four vertical sections, one for each of the synthesized sounds, and designed to resemble a kitchen top - an oven, specifically. The "Play" starts the selected sound playback, while the sliders allow the user to control the volume and the panning of the played sound. The square corresponding to the oven door is the plot of the stereo

sound (left and right channel) either in the time domain or in the frequency domain. The yellow button is deputed to switching the domain of representation of the sound. The two rows of knobs control the sound parameters that are customizable by the user: the first row manages the granular synthesis parameters, while the second contains the reverb controls.



Figure 5. Application GUI.

4. Conclusions

What became clear during the process is that only by means of a deeper understanding of granular synthesis, achievable only through experience and time, the sounds generated with this technique could have been more accurate. Understanding how to modify the granulator parameters, during the first phases of the work, in order to cause certain effects on the sound features, was not at all easy. It was therefore difficult to choose how many and which sounds it was possible to reproduce, and which parameters could be customized by the user without dramatically changing the character of the sound.

Nevertheless, the authors feel confident about the results, in terms of the quality of the sounds and the usability of the application.